

# Fraud Detection System for Banking Transactions

## Comprehensive Documentation for Word Copy-Paste

### 1. Algorithm Justification: What, Why, and Suitability

#### Overview of Chosen Algorithms

- **K-means Clustering (Unsupervised)**
- **Random Forest (Supervised)**
- **Logistic Regression (Supervised)**
- **Gradient Boosted Trees (Supervised)**

#### Why Only These Algorithms?

- **K-means Clustering:**
  - Identifies outliers and patterns in data without requiring fraud labels.
  - Enables the discovery of new, previously unseen types of fraud.
  - Works efficiently with large datasets in Apache Spark, making it scalable for banking data.
- **Random Forest:**
  - Offers high accuracy and resilience against overfitting due to ensembling multiple decision trees.
  - Handles mixed types of variables and high-dimensional data, making it ideal for diverse banking transaction features.
  - Supports feature importance analysis, helping explain predictions to compliance teams.
- **Logistic Regression:**
  - Provides clear, interpretable probabilistic outputs—essential for regulatory compliance and explaining model behavior to stakeholders.
  - Fast training and inference, especially valuable for very large, sparse datasets.
  - Suitable as a baseline model for benchmarking more complex approaches.

- **Gradient Boosted Trees:**

- Excels at detecting complex, nonlinear fraud patterns by sequentially combining weak learners.
- Typically achieves very high predictive performance (precision, recall, AUC) on fraud datasets.
- Well-supported and optimized in Spark MLlib for scalable model training.

## **Suitability for This Project**

- All algorithms are available and optimized for distributed processing with Apache Spark, crucial for handling millions of transactions.
- The combination provides a full spectrum for both **anomaly detection** (K-means) and **fraud classification** (other three models).
- Real-time model scoring is possible via Spark Streaming and distributed ML pipelines.
- These models lend themselves well to feature importance analysis, threshold tuning, and integration with real-world production alerting systems.

## **2. Source Code Reference (Overview & Sample)**

### **Architecture & Implementation Outline**

- **Data Ingestion**
  - Spark or Hadoop reads data sources in parallel for spatial and temporal scalability.
  - Transaction features: date, code, amount, and engineered features (e.g., time since last transaction).
- **Preprocessing & Feature Engineering**
  - Handling missing/null values and standardizing formats.
  - Feature scaling and encoding (using VectorAssembler, StandardScaler in Spark).
  - Generating customer-centric features (frequency, transaction variance).
- **Model Training**
  - Modular scripts for model selection, training, and hyperparameter tuning.
  - K-means handles unsupervised clustering and assigns "anomaly" scores.

- Supervised models (Random Forest, Logistic Regression, Gradient Boosted Trees) fit labels and predict fraud probabilities.
- **Evaluation**
  - Extensive metrics calculation: AUC, precision, recall, F1-score, confusion matrix.
  - Cross-model comparison via ROC curves.
  - Feature importance charts for interpretability.
- **Deployment & Inference**
  - Spark MLlib pipelines for quick retraining and inference.
  - Spark Streaming for live scoring in production.

### Sample Spark ML Pipeline Snippet

```
from pyspark.ml.feature import VectorAssembler, StandardScaler
from pyspark.ml.clustering import KMeans
from pyspark.ml.classification import RandomForestClassifier, LogisticRegression, GBTClassifier
from pyspark.ml import Pipeline

# 1. Feature Assembling and Scaling
assembler = VectorAssembler(inputCols=["amount", "hour", "day"], outputCol="features")
scaler = StandardScaler(inputCol="features", outputCol="scaledFeatures")

# 2. Model Definitions
kmeans = KMeans(featuresCol="scaledFeatures", k=5)
rf = RandomForestClassifier(featuresCol="scaledFeatures", labelCol="label", numTrees=100)
lr = LogisticRegression(featuresCol="scaledFeatures", labelCol="label")
gbt = GBTClassifier(featuresCol="scaledFeatures", labelCol="label")

# 3. Training Pipeline (example with Random Forest)
pipeline_rf = Pipeline(stages=[assembler, scaler, rf])
model_rf = pipeline_rf.fit(trainingData)
```

```
predictions_rf = model_rf.transform(testData)
```

**Note:** Replace ["amount", "hour", "day"] with your actual feature names as per real dataset.

- For full source code, refer to the main GitHub repository:  
[\[github.com/karunakar00/Fraud-Detection-in-Banking-Transactions-Using-Hadoop/tree/main\]](https://github.com/karunakar00/Fraud-Detection-in-Banking-Transactions-Using-Hadoop/tree/main)

### 3. Comprehensive Project Documentation

#### Project Pipeline

##### Step 1: Data Loading & Exploratory Analysis

- Use Spark to ingest large-scale transaction data.
- Perform EDA to understand fraud rates, feature distributions, and temporal trends.
- Visualize data with histograms, scatter plots, and correlation matrices.

##### Step 2: Data Cleaning & Feature Engineering

- Remove or impute missing values.
- Standardize numeric features needed for K-means and regression.
- Engineer new features: transaction frequency, average amount per user, ratios, and time-based features.

##### Step 3: Model Development

| Algorithm              | Type         | Core Strengths                    | Fraud Detection Role                      |
|------------------------|--------------|-----------------------------------|---|
| K-means Clustering     | Unsupervised | Fast, scalable. Outlier detection | Identifies novel/suspicious transactions  |
| Random Forest          | Supervised   | Robust, accurate, interpretable   | Main classifier for binary fraud labels   |
| Logistic Regression    | Supervised   | Simple, fast, interpretable       | Baseline, regulatory explanations         |
| Gradient Boosted Trees | Supervised   | Top AUC, nonlinear patterns       | Captures complex, subtle fraud signatures |

## Step 4: Model Evaluation & Comparison

- Use cross-validation for benchmarking.
- Calculate precision, recall, F1-score, and plot ROC curves for each model.
- Interpret feature importances—highlighting which attributes most influence the fraud decision.

## Step 5: Real-time Prediction & Deployment

- Integrate with Kafka or Spark Streaming for instant fraud detection.
- Send real-time alerts for flagged transactions to downstream systems (e.g., customer notification, security review).
- Monitor model performance and retrain regularly to address emerging fraud tactics.

## 4. Detailed Graphs & Visual Insights

- **Feature Importance Bar Chart (Random Forest):**  
Visualizes the top transaction features contributing to fraud prediction (e.g., amount, transaction time, transaction type).
- **ROC Curve Comparison:**  
Plots ROC curves for all algorithms, illustrating how well each model distinguishes fraudulent from legitimate transactions.
- **Cluster Scatter Plot (K-means):**  
Shows transaction clusters; outliers (far from centroids) may correspond to potential fraud.

## 5. Deployment Recommendations

- **Environment:**  
Running on Google Colab, Databricks, or any Spark/Hadoop-supported cluster.
- **Steps:**
  - a. Clone the referenced repository and load the main pipeline scripts.
  - b. Replace the sample dataset with real transaction data (e.g., Kaggle's `creditcard.csv`).
  - c. Adjust feature engineering and pipeline stages as needed.
  - d. Run the training and evaluation cells.

- e. For production, enable Spark Streaming and connect to Kafka for ingesting live transactions.

## 6. References

- Main Source:  
[github.com/karunakar00/Fraud-Detection-in-Banking-Transactions-Using-Hadoop/tree/main](https://github.com/karunakar00/Fraud-Detection-in-Banking-Transactions-Using-Hadoop/tree/main)
- Example Datasets:
  - Credit Card Fraud Detection, Kaggle
  - Synthetic Financial Datasets For Fraud Detection, Kaggle

**Presented by:-**

**Name:- PONNADA KUMARASWAMY**

**ROLL N.O:- 24M11MC266**

**COLLEGE:- ADITYA UNIVERSITY**